

Taking Global Snapshots of a Distributed Banking Application

1 A Distributed Banking Application

In the first part, you will implement a distributed banking application. The distributed bank has multiple branches. Every branch knows about all other branches. Each branch starts with some initial balance. It then randomly selects a branch and send a random amount of money to this branch at unpredictable times. You must implement the following functions for a branch in a distributed bank:

`initBranch` this method has two input parameters: the initial balance of a branch and a list of all branches in the distributed bank. Upon receiving this method, a branch will set its initial balance and record the list of all branches.

`transferMoney` given a `TransferMessage` structure that contains the sending branch's `BranchID`, the amount of money in the message, as well as a `messageId` (more on this later), the receiving branch updates its balance accordingly.

A branch can be both a sender and a receiver. Therefore, you need to implement a method that will be called from a separate thread in your branch to perform the sending operation. Intervals between consecutive sending operations should be drawn uniformly at random between 0 and 5 seconds. A sender can only send positive amount of money. It needs to first decrease its balance, then call the `transferMoney` method on a remote branch. A branch's balance should not become negative. For simplicity, the amount of money should be drawn randomly between 1% and 5% of the branch's initial balance and can only be an integer.

The branch's main program should spawn both a *multi-threaded Thrift server* and also the *thread that sends money from your branch to other branches*. Because the balance of your branch will be read and written by multiple threads, this variable needs to be protected by a mutex or another synchronization method. In addition, you can assume that neither the branches nor the communication channels will fail. So you don't need to implement anything to guarantee fault tolerant transactions.

Your branch executable should take two command line inputs. The first one is a human-readable name of the branch, e.g., "branch1". The second one specifies the port number the branch runs on.

```
$> ./branch branch1 9090
```

1.1 Controller

In this project, we rely on a controller to set a branch's initial balance and notify every branch of all other branches in the distributed bank. This controller takes two command line inputs: the total amount of money in the distributed bank and a local file that stores the names, IP addresses, and port numbers of all branches.

An example of how the controller program should operate is provided below:

```
$> ./controller 4000 branches.txt
```

The file (`branches.txt`) should contain a list of names, IP addresses, and ports, in the format "<name> <public-ip-address> <port>", of all of the running branches.

For example, if four branches with names: “branch1”, “branch2”, “branch3”, and “branch4” are running on remote01.cs.binghamton.edu port 9090, 9091, 9092, and 9093, then branches.txt should contain:

```
branch1 128.226.180.163 9090 branch2  
128.226.180.163 9091 branch3 128.226.180.163  
9092 branch4 128.226.180.163 9093
```

The controller will distribute the total amount of money evenly among all branches, e.g., in the example above, every branch will receive \$1,000 initial balance. The controller initiates all branches by individually calling the `initBranch` method described above. Note that the initial balance must be integer.

2 Taking Global Snapshots of the Bank

In the second part, you will use the Chandy-Lamport global snapshot algorithm take global snapshots of your bank. In case of the distributed bank, a global snapshot will consist of the local state of each branch (i.e., its balance) and the amount of money in transit on all communication channels. Each branch will be responsible for recording and reporting its own local state (balance) as well as the total money in transit on each of its incoming channels.

For simplicity, in this project, the controller will contact one of the branches to initiate the global snapshot. It does so by making a `initSnapshot` call to the selected branch. The selected branch will then initiate the snapshot by first recording its own local state and send out Marker messages to all other branches. After some time (long enough for the snapshot algorithm to finish), the controller makes `retrieveSnapshot` calls to all branches to retrieve their recorded local and channel states. Note that if the snapshot is correct, the total amount of money in all branches and in transit should equal to the command line argument given to the controller. You need to implement the following three functions to add snapshot capability to your distributed bank.

`initSnapshot` upon receiving this call, a branch records its own local state (balance) and sends out Marker messages to all other branches by calling the `Marker` method on them. To identify multiple snapshots, the controller passes in a `snapshotId` to this call, and all the marker messages should include this `snapshotId`.

`Marker` given the sending branch’s `BranchID`, `snapshotId`, and `messageld`, the receiving branch does the following:

1. if this is the first Marker message with the `snapshotId`, the receiving branch records its own local state (balance), records the state of the incoming channel from the sender to itself as empty, immediately starts recording on other incoming channels, and sends out Marker messages to all of its outgoing channels (i.e., all branches except itself).
2. otherwise, the receiving branch records the state of the incoming channel as the sequence of money transfers that arrived between when it recorded its local state and when it received the Marker. `retrieveSnapshot` given the `snapshotId` that uniquely identifies a snapshot, a branch retrieves its recorded local and channel states and return them to the caller (i.e., the controller).

The controller should be fully automated. It periodically calls the `initSnapshot` method with monotonically increasing `snapshotId` on a randomly selected branch and outputs to the screen the aggregated global snapshot retrieved from all branches in a human-readable manner. In addition, the snapshot taken by branches needs to

be identified by their names: e.g., “branch1” to represent branch1’s local state, and “branch2->branch1” to represent the channel state.

2.1 FIFO message delivery

The correctness of the Chandy-Lamport snapshot algorithm relies on FIFO message delivery of all communication channels among all branches (processes). A communication channel is a one way connection between two branches. For example, in this project, from “branch1” to “branch2” is one communication channel. From “branch2” to “branch1” is another channel. In order to ensure FIFO message delivery, we can include a monotonically increasing messageId in all messages transferred in a channel.

For instance, suppose that “branch2” receives a transferMoney message or a Marker message from “branch1” with messageId 5, but the last messageId that it observed from “branch1” was 3. In this case, “branch2” would be able to infer that it is missing the message with messageId 4 from “branch1”. When “branch2” sees that it is missing a message, it will pause its processing of the message with messageId 5 and wait for the message with messageId 4 to first arrive and then be processed.

The sender branch must maintain a last used messageId for each of its outgoing communication channels. In this project, two types of messages are transferred among branches (processes): transferMoney and Marker messages. Before sending a transferMoney or Marker message, the sender must increment the last used messageId of appropriate outgoing channel and include this messageId in the request. In addition, since each branch sends transferMoney and Marker messages in different threads, you must also use a mutex to protect each last used messageId variable for each outgoing communication channel.

For each incoming channel, a last seen messageId should be maintained at each branch. On receipt of a new message, you can use a condition variable to wait for all prior messages of this incoming channel with lower messageId values to be processed. After completing message processing on the channel, you should increment the last seen messageId.